

Shou Hao (Lennox) Ho

Email:// lennoxhoe@gmail.com LinkedIn:// ["Lennox" Shou Hao Ho](#)

Personal Github:// [lennoxho](#) Work Github:// [holenno1](#)

Tel: (+1)416-788-4307

WORK EXPERIENCE

INTEL PSG (FORMERLY ALTERA) | Senior Software Developer, Grade 7

June 2016 - Current | Toronto, ON Canada

- Software developer for the [Intel Quartus Prime](#) software
- Quartus **database and design flows** team. Responsibilities include:
 - Development of custom **data structures** for netlists and ancilliary data
 - Ensure correct and efficient **exchange of data** between Quartus modules
 - **Orchestration** of Quartus processes throughout a flow
- Maintainer of Quartus **common code infrastructure**. Responsibilities include:
 - Development of **core in-house libraries** and integration of **third party libraries**
 - Development of **internal debugging and profiling tools**
 - Planning and main driver of **compiler upgrades**
 - Maintain the **overall stability and health** of the Intel Quartus Prime software
 - De facto **engineering representative** for infrastructure & OS support discussions

ACHIEVEMENTS

- Fixed a **critical defect** blocking Quartus **support on Windows 11**. Departmental Recognition Award (DRA). [More...](#)
- Complete rewrite of Quartus in-house **memory profiler** (qmp), **runtime overhead** reduced by **10x**, **memory overhead** by **50%**. [More...](#)
- Clean room implementaton of **DWARF** debugging format (v2-v5) **decoder**. Tailored to [qmp](#)'s runtime and memory requirements. [More...](#)
- Created an **Electron**-based **GUI** for [qmp](#). Power tool to analyse memory usage across a process lifecycle. [More...](#)
- Reduced **build time** of Quartus binaries by **400%**. Integration test runtime reduced by 3 hours
- Reduced Quartus **devkit setup time** on Windows by **600%** (1 hour to 10 minutes)
- Fixed performance of **TCMalloc** on Windows, yielding a **10% global reduction** in Quartus **runtime**. [More...](#)
- Complete rewrite of Quartus in-house **database schema parser/generator**, leveraging **libclang** to parse C++ code. [More...](#)
- Rewrote Quartus **wildcard assignment processing** engine, achieving a **90% speedup**. [More...](#)
- Migrated Quartus netlist databases from **memory mapping to heap memory**, yielding a **50%** reduction in **virtual memory** usage
- Spearheaded **ABI and runtime compatibility** checking for Quartus as an engineering-driven approach of ensuring product stability on a wide range of platforms
- Designed a **novel quasi-succinct fixed-data hash table** to store associative data in Quartus device databases. [More...](#)
- Designed a **novel mutex methodology** using thread local storage, used to implement **semi-lock-free data structures** that are significantly more scalable than contemporaries. [More...](#)

PROGRAMMING LANGUAGES | SKILLS

PROFICIENT

- C++ 20 • C • Python
- Microsoft Windows • Linux
- Concurrent & parallel programming • Lock-free programming
- Systems programming • Low-level optimisations • Synchronisation primitives
- Template Metaprogramming • Code & assembly-level debugging

FAMILIAR

- D • Rust • Javascript • Java
- x86 assembly • x86 SIMD programming

TOOLS

- Microsoft Visual Studio • gdb • Boost • libClang • Intel TBB • libunwind • CMake • Electron • Node-API

EDUCATION

UNIVERSITY OF TORONTO | BSc IN COMPUTER SCIENCE (HONS.)

Specialising in Scientific Computing | CGPA 3.95/4.0

LEN'S HOT TAKES

- Weekly blog series on Intel PSG's internal forums, doing deep dives into obscure topics such as the implementation of thread local storage, debugging runtime issues etc
- "-Bsymbolic and std::string" explores how a module built with -Bsymbolic and the old libstdc++ ABI (uses copy-on-write strings) causes weird crashes
- "The mystery of the __std_type_info_name crash" explores how using an 8-byte aligned allocator results in an instruction-level fault
- Since December 2020

PROJECTS

QUARTUS WINDOWS 11 COMPATIBILITY

- A [critical defect](#) was found late in the development cycle affecting all **Quartus GUI applications** on **Windows 11**
- Used [trampoline hooking](#) in **TCMalloc** to work around problematic Windows runtime behaviour
- Worked with Intel TBB maintainers to evaluate possible mitigations in **TBBMalloc**. Implemented [import address table \(IAT\) hooking](#) in TBBMalloc
- Recognised with a **Departmental Recognition Award (DRA)**

TCMALLOC PERFORMANCE ON WINDOWS

- [TCMalloc](#) is a high performance allocation library developed by Google
- Quartus had been using TCMalloc on Linux for a number of releases
- TBBMalloc was used on Windows, as TCMalloc appeared to offer lower performance
- **Identified the reason for TCMalloc slowdown on Windows** and implemented a fix, improving TCMalloc performance on Windows by **10%**
- [Upstreamed](#) fixes and enhancements to the TCMalloc public repository

QUARTUS MEMORY PROFILER (qmp) V2

- Rewrote the Quartus in-house **memory profiling tool** from first principles as a completely standalone tool
- Achieved **1000% reduction in runtime overhead** while **memory overhead is reduced by 50%**
- Typical runtime overhead of memory profiling is now around 35% compared to no profiling
- This speedup was made possible using heavily modified copies of **libunwind** and **TCMalloc**, as well as custom **lock-free data structures** ensuring wait-free operation
- Enabled memory profiling to be successfully performed on the entire **Quartus competitive benchmarking design suite** with a **24 hour turnaround time**
- Features custom variable-length encoding format and data deduplication efforts to minimise size of results (\approx 50 MBs / hour of collection)

DWARF DEBUGGING FORMAT DECODER

- Clean room implementation of [DWARF](#) debugging format decoder with no 3rd party dependency
- **Versions 2-5** (most recent as of writing) of the debugging format are supported with the exception of split and supplementary-DWARF
- Caches decoded data in memory for fast lookup from [qmp](#) while minimising memory overhead
- Also exposed `qmp-addr2line` as a GNU [addr2line](#) clone

QMP GUI

- Designed and implemented an **Electron**-based GUI to display [qmp](#) profiling results
- Custom **Node-API**-based backend to perform data marshalling and heavy number crunching in C++
- **Shared** memory used in backend to avoid duplication of profile data across processes
- Heavily customised popular Javascript data visualisation libraries for real time performance on GBs of compressed data
- Implemented a **Level-of-detail** (LOD) system for the [Dygraphs](#) library to exceed the recommended 10,000 data point limit
- Features a timeline (stacked) chart of memory usage broken down by module over the profiling duration
- A Flame chart representing the call stack is overlayed and synchronised with the timeline (memory) chart. This allows Quartus developers to track the lifecycle of any byte of memory
- Fully interactive, including zooming and panning support

QUARTUS DATABASE SCHEMA PARSER V2

- Quartus' in-house **serialisation library** ships with a tool to automatically **generate serialisation routines (schemas)** based on C++ code
- This parser/generator tool featured a handwritten C++ parser, which was fragile and frequently broke whenever a new compiler was introduced
- **Rewrote the parser/generator tool** to leverage [libclang](#) as its C++ parser
- This change improved the accuracy of the generated schemas while ensuring the long term viability of the library

QUARTUS WILDCARD ASSIGNMENT PROCESSING

- Rewrote Quartus' **handling of wildcard assignments**
- Achieved a **90% speedup** by leveraging the **hierarchical nature of netlists** and how that relates to the **"greedy"** nature of **simple wildcard patterns**

QUASI-SUCCINCT FIXED-DATA HASH TABLE

- Designed and implemented a **novel quasi-succinct hash table** data structure tailored for **precomputed data**
- Open-addressing & linear-probing based, where the worst case probe length was the size of the largest bucket

- **In memory compression** (inspired by **Elias-Fano encoding**) was used to store hash table indices. **x86 SIMD intrinsics** and jump tables were used to perform fast bucket offset lookup
- Invented as a solution to **eliminate wastage** when traditional associative data structures were used for **fixed data**

HYBRID SHARED MUTEX & SEMI-LOCK FREE DATA STRUCTURES

- Designed to overcome the traditional limitations of conventional lock-free data structures, chiefly the need for individually-allocated nodes
- The node-based design is necessary for lock-free techniques to be effectively deployed
- While various optimisations exist, node-based data structures generally have a high memory overhead
- **Hybrid shared mutex** leverages **thread-local-storage** to provide very **cheap read locks**, and more **expensive write locks**
- **Semi-lock free data structures** may use read locks to perform atomic operations on a **fixed size bufer**, while a write lock is used to increase the size of the buffer
- With a well chosen growth factor, the **vast majority of operations** on a semi-lock free data structure will be **lock-free**
- The ability to use **fixed size bufers** in semi-lock free data structures allows for more **compact implementations of high performance, thread-safe data structures**

QUARTUS TCL VIRTUAL FILESYSTEM

- Implemented a **virtual filesystem layer** to provide the **Quartus Tcl interpreter** direct access to Quartus databases
- This greatly improves the project migration (export/import) experience by packaging all **user SDC and Tcl scripts** directly into the **Quartus databases**

CLANG ONBOARDING FOR QUARTUS

- Incrementally **fixed the** (\approx 30 million LOC) **Quartus codebase** to be **compatible with the Clang compiler**
- Clang compliance is a **prerequisite** for a number of **Clang-based static analysis tools** such as SourceTrail and ClangTidy
- This also enables Quartus to be built with the [Clang compiler](#) or the [Intel ICX compiler](#)

COMPILER AND C++ STANDARD UPDATES FOR QUARTUS

- Led efforts to update Quartus to MSVC 2019 and 2022, GCC 12 and Clang 15
- Led efforts to upgrade the Quartus codebase to **C++14, 17, then 20**