# Lennox Shou-Hao Ho

Email:// lennoxhoe@gmail.com  LinkedIn:// "Lennox" Shou Hao Ho

Personal Github:// lennoxho  Tel: (+1)416-788-4307

## WORK EXPERIENCE

### EFINIX | Software Developer, SMTS

Feb 2024 - Current | Toronto, ON Canada

- Software developer for the Efinity EDA software

- Project lead for **flows, databases and code Infrastructure** (FDI). Responsibilities include:
  - Design and implementation of high level **compiler control flows**
  - Profiling, monitoring and pathfinding for **runtime, memory and load time** improvements
  - Setup and maintenance of local **SLURM compute farm**

### INTEL PSG / ALTERA | Software Developer, Grade 7

June 2016 - Dec 2023 | Toronto, ON Canada

- Software developer for the Intel Quartus Prime EDA software

- Co-owner of Quartus **common code infrastructure**. Responsibilities include:
  - Design and implementation of **internal debugging and profiling tools** and **core in-house libraries**
  - De facto **engineering representative** for infrastructure & OS support discussions

- Maintainer of Quartus **database and design flows**. Responsibilities include:
  - Design and implementation of custom **data structures** for netlists and ancilliary data

## ACHIEVEMENTS

- In the first year at Efinix,
  - Reduced Efinity **synthesis runtime by 70%** through a series of micro-optimisations, algorithm reformulation to enable parallelism
  - Reduced Efinity **place & route runtime by 40%** through a series of micro-optimisations, data reorganisation for better locality
  - Reduced Efinity **memory usage by 3x**. Dropped Efinity memory requirements from **32GB to 12GB (16 threads)**
  - Reduced Efinity **load time by 5x**, from up to **20 minutes down to 3 minutes** on largest designs
  - Replaced the **Efinity SDC parser** with a compliant Tcl interpreter
  - **Trivialised wildcard processing time** in SDC scripts by implementing an SSE-optimised wildcard matching engine
  - Set up a local SLURM **compute farm** to support more QoR experiments

- During tenure at Altera,
  - Fixed a **system library defect** blocking Quartus support on **Windows 11**. Departmental Recognition Award (DRA).
  - First-principles rewrite of Quartus in-house **memory profiler** (qmp). Profiling **overhead reduced by 10x** compared to previous effort
  - Rewrote Quartus in-house **database parser/generator/serialisation library** by leveraging libClang to parse C++ code
  - Clean room implementaton of DWARF debugging format (v2-5) decoder. Tailored to **qmp**'s runtime and memory requirements
  - Designed and implemented an Electron-based **GUI** for **qmp**. Power tool to analyse memory usage across a process lifecycle
  - Fixed performance of open-source TCMalloc library on Windows, yielding a **10% global reduction** in Quartus **runtime**
  - Reduced **build time** of Quartus binaries by **4x (3 hours)**, **devkit setup time** on Windows by **6x (50 minutes)**
  - Designed a **novel quasi-succinct fixed-data hash table** to store prebuilt and constant associative data in Quartus device databases
  - Designed a **novel mutex methodology** using thread local storage, used to implement **semi-lock-free data structures**

## PROGRAMMING LANGUAGES | SKILLS

### PROFICIENT

- **C++ 23** • **C** • **Python**       • Concurrent & parallel programming • Lock-free programming • Synchronisation primitives
- **Microsoft Windows** • **Linux**    • Template Metaprogramming • x86 SIMD programming • Instruction-level optimisation

### FAMILIAR

- **D** • **Rust** • **Javascript** • **Java**   • x86 assembly

### TOOLS

- MSVC • gdb • libClang • LLVM • Boost • Intel TBB • libunwind • CMake • Electron • Node-API

## EDUCATION

**UNIVERSITY OF TORONTO** | BSc in Computer Science (Hons.)

Specialising in Scientific Computing | CGPA 3.95/4.0

## PROJECTS

### QUARTUS WINDOWS 11 COMPATIBILITY

- Critical defect found late in the development cycle affected all **Quartus GUI applications** on **Windows 11**

- Used trampoline hooking in **TCMalloc** to work around problematic Windows runtime behaviour

- Worked with Intel TBB maintainers to evaluate possible mitigations in **TBBMalloc**. Implemented import address table (IAT) hooking in TBBMalloc

- Recognised with a **Departmental Recognition Award (DRA)**

### TCMALLOC PERFORMANCE ON WINDOWS

- TCMalloc is a high performance allocation library developed by Google

- TBBMalloc was used on the Windows build of Quartus, as TCMalloc appeared to offer lower performance

- **Identified the reason for TCMalloc slowdown on Windows** and implemented a fix, improving TCMalloc performance on Windows by **10%**

- Upstreamed fixes and enhancements to the TCMalloc public repository

### QUARTUS MEMORY PROFILER (qmp) V2

- Redesigned the Quartus in-house **memory profiling tool** from first principles as a completely standalone tool

- Achieved **10x reduction in runtime overhead** while **memory overhead is reduced by 2x**

- Typical runtime overhead of memory profiling is reduced to 35% compared to no profiling

- This speedup was made possible using heavily modified copies of **libunwind** and **TCMalloc**, as well as custom **lock-free data structures** ensuring wait-free operation

- Enabled memory profiling to be successfully performed on the entire **Quartus competitive benchmarking design suite** with a **24 hour turnaround time**

- Features custom variable-length encoding format and data deduplication efforts to minimise size of results ($\approx 50$ MBs / hour of collection)

### DWARF DEBUGGING FORMAT DECODER

- Clean room implementation of DWARF debugging format decoder with no $3^{rd}$ party dependency

- **Versions 2-5** (most recent as of writing) of the debugging format are supported with the exception of split and supplementary-DWARF

- Also exposed `qmp-addr2line` as a GNU `addr2line` clone

## QMP GUI

• Designed and implemented an **Electron**-based GUI to display **qmp** profiling results

• Custom **Node-API**-based backend to perform data marshalling and heavy number crunching in C++

• **Shared** memory used in backend to avoid duplication of profile data across processes

• Heavily customised popular Javascript data visualisation libraries for real time performance on GBs of compressed data

• Implemented a **Level-of-detail** (LOD) system for the **Dygraphs** library to exceed the recommended 10, 000 data point limit

• Features a timeline (stacked) chart of memory usage broken down by module over the profiling duration

• A Flame chart representing the call stack is overlayed and synchronised with the timeline (memory) chart. This allows Quartus developers to track the lifecycle of any byte of memory

• Fully interactive, including zooming and panning support

## FLOWGRAPH-DRIVEN PARALLEL SERIALISATION LIBRARY

• Trivialised Quartus netlist database loading time by designing and implementing a **flowgraph-based serialisation library**

• A flowgraph is used to represent the Quartus **data dependency graph**, which when fed to a parallel computing library can yield huge reductions in loading time

## QUARTUS DATABASE SCHEMA PARSER V2

• Quartus' in-house **serialisation library** ships with a tool to automatically **generate serialisation routines (schemas)** based on C++ code

• This parser/generator tool featured a handwritten C++ parser, which was fragile and frequently broke whenever a new compiler was introduced

• **Redesigned and rewrote the parser/generator tool** to leverage **libclang** as its C++ parser

• This change improved the accuracy of the generated schemas while ensuring the long term viability of the library

## QUASI-SUCCINCT FIXED-DATA HASH TABLE

• Designed and implemented a **novel quasi-succint hash table** data structure tailored for **precomputed data**

• Open-addressing & linear-probing based, where the worst case probe length was the size of the largest bucket

• **In memory compression** (inspired by **Elias-Fano encoding**) was used to store hash table indices. **x86 SIMD instrinsics** and jump tables were used to perform fast bucket offset lookup

• Invented as a solution to **eliminate wastage** when traditional associative data structures were used for **fixed data**

## HYBRID SHARED MUTEX & SEMI-LOCK FREE DATA STRUCTURES

• Designed to overcome the traditional limitations of conventional lock-free data structures, chiefly the need for individually-allocated nodes

• While various optimisations exist, node-based data structures generally have a high memory overhead

• **Hybrid shared mutex** leverages **thread-local-storage** to provide very **cheap read locks** due to its hardware cache advantages over shared atomics

• **Semi-lock free data structures** may use read locks to perform atomic operations on a **fixed size bufer**, while a write lock is used to increase the size of the buffer

• The ability to use **fixed size bufers** in semi-lock free data structures allows for more **compact implementations of high performance, thread-safe data structures**

## QUARTUS TCL VIRTUAL FILESYSTEM

• Implemented a **virtual filesystem layer** to provide the **Quartus Tcl interpreter** direct access to Quartus databases

• This greatly improves the project migration (export/import) experience by packaging all **user SDC and Tcl scripts** directly into the **Quartus databases**