

# Lennox Shou-Hao Ho

Email:// [lennoxhoe@gmail.com](mailto:lennoxhoe@gmail.com) LinkedIn:// "Lennox" Shou Hao Ho

Personal Github:// [lennoxho](https://github.com/lennoxho) Tel: (+1)416-788-4307

## WORK EXPERIENCE

### EFINIX | Software Developer, MTS

Feb 2024 - Current | Toronto, ON Canada

- Software developer for the [Efinity](#) EDA software
- Project lead for **flows, databases and code Infrastructure** (FDI). Responsibilities include:
  - Design and implementation of high level **compiler control flows**
  - Profiling, monitoring and pathfinding for **runtime, memory and load time** improvements
  - Setup and maintenance of local **SLURM** compute farm

### INTEL PSG / ALTERA | Software Developer, Grade 7

June 2016 - Dec 2023 | Toronto, ON Canada

- Software developer for the [Intel Quartus Prime](#) EDA software
- Co-owner of Quartus **common code infrastructure**. Responsibilities include:
  - Design and implementation of **internal debugging and profiling tools** and **core in-house libraries**
  - De facto **engineering representative** for infrastructure & OS support discussions
- Maintainer of Quartus **database and design flows**. Responsibilities include:
  - Design and implementation of custom **data structures** for netlists and ancillary data

## ACHIEVEMENTS

- In the first year at Efinix,
  - Reduced Efinity **synthesis runtime by 70%** through a series of micro-optimisations, algorithm reformulation to enable parallelism
  - Reduced Efinity **place & route runtime by 40%** through a series of micro-optimisations, data reorganisation for better locality
  - Reduced Efinity **memory usage by 3x**. Dropped Efinity memory requirements from **32GB to 12GB (16 threads)**
  - Reduced Efinity **load time by 5x**, from up to **20 minutes down to 3 minutes** on largest designs
  - Replaced the **Efinity SDC parser** with a compliant Tcl interpreter
  - **Trivialised wildcard processing time** in SDC scripts by implementing an **SSE**-optimised wildcard matching engine
  - Set up a local **SLURM** compute farm to support more QoR experiments
- During tenure at Altera,
  - Fixed a **system library defect** blocking Quartus support on **Windows 11**. Departmental Recognition Award (DRA).
  - First-principles rewrite of Quartus in-house **memory profiler** (qmp). Profiling **overhead reduced by 10x** compared to previous effort
  - Rewrote Quartus in-house **database parser/generator/serialisation library** by leveraging **libClang** to parse C++ code
  - Clean room implementaton of **DWARF debugging format (v2-5) decoder**. Tailored to **qmp**'s runtime and memory requirements
  - Designed and implemented an **Electron**-based **GUI** for **qmp**. Power tool to analyse memory usage across a process lifecycle
  - Fixed performance of open-source **TCMalloc** library on Windows, yielding a **10% global reduction** in Quartus **runtime**
  - Reduced **build time** of Quartus binaries by **4x (3 hours)**, **devkit setup time** on Windows by **6x (50 minutes)**
  - Designed a **novel quasi-succinct fixed-data hash table** to store prebuilt and constant associative data in Quartus device databases
  - Designed a **novel mutex methodology** using thread local storage, used to implement **semi-lock-free data structures**

## PROGRAMMING LANGUAGES | SKILLS

### PROFICIENT

- C++ 23 • C • Python
- Concurrent & parallel programming • Lock-free programming • Synchronisation primitives
- Microsoft Windows • Linux
- Template Metaprogramming • x86 SIMD programming • Instruction-level optimisation

### FAMILIAR

- D • Rust • Javascript • Java
- x86 assembly

### TOOLS

- MSVC • gdb • libClang • LLVM • Boost • Intel TBB • libunwind • CMake • Electron • Node-API

# EDUCATION

UNIVERSITY OF TORONTO | BSc IN COMPUTER SCIENCE (HONS.)

Specialising in Scientific Computing | CGPA 3.95/4.0

## PROJECTS

### QUARTUS WINDOWS 11 COMPATIBILITY

- **Critical defect** found late in the development cycle affected all **Quartus GUI applications** on **Windows 11**
- Used **trampoline hooking** in **TCMalloc** to work around problematic Windows runtime behaviour
- Worked with Intel TBB maintainers to evaluate possible mitigations in **TBBMalloc**. Implemented **import address table (IAT) hooking** in TBBMalloc
- Recognised with a **Departmental Recognition Award (DRA)**

### TCMALLOC PERFORMANCE ON WINDOWS

- **TCMalloc** is a high performance allocation library developed by Google
- TBBMalloc was used on the Windows build of Quartus, as TCMalloc appeared to offer lower performance
- **Identified the reason for TCMalloc slowdown on Windows** and implemented a fix, improving TCMalloc performance on Windows by **10%**
- **Upstreamed** fixes and enhancements to the TCMalloc public repository

### QUARTUS MEMORY PROFILER (qmp) V2

- Redesigned the Quartus in-house **memory profiling tool** from first principles as a completely standalone tool
- Achieved **10x reduction in runtime overhead** while **memory overhead is reduced by 2x**
- Typical runtime overhead of memory profiling is reduced to 35% compared to no profiling
- This speedup was made possible using heavily modified copies of **libunwind** and **TCMalloc**, as well as custom **lock-free data structures** ensuring wait-free operation
- Enabled memory profiling to be successfully performed on the entire **Quartus competitive benchmarking design suite** with a **24 hour turnaround time**
- Features custom variable-length encoding format and data deduplication efforts to minimise size of results ( $\approx 50$  MBs / hour of collection)

### DWARF DEBUGGING FORMAT DECODER

- Clean room implementation of **DWARF** debugging format decoder with no **3<sup>rd</sup>** party dependency
- **Versions 2-5** (most recent as of writing) of the debugging format are supported with the exception of split and supplementary-DWARF
- Also exposed **qmp-addr2line** as a GNU **addr2line** clone

## QMP GUI

- Designed and implemented an **Electron**-based GUI to display **qmp** profiling results
- Custom **Node-API**-based backend to perform data marshalling and heavy number crunching in C++
- **Shared** memory used in backend to avoid duplication of profile data across processes
- Heavily customised popular Javascript data visualisation libraries for real time performance on GBs of compressed data
- Implemented a **Level-of-detail** (LOD) system for the **Dygraphs** library to exceed the recommended 10,000 data point limit
- Features a timeline (stacked) chart of memory usage broken down by module over the profiling duration
- A Flame chart representing the call stack is overlaid and synchronised with the timeline (memory) chart. This allows Quartus developers to track the lifecycle of any byte of memory
- Fully interactive, including zooming and panning support

## FLOWGRAPH-DRIVEN PARALLEL SERIALISATION LIBRARY

- Trivialised Quartus netlist database loading time by designing and implementing a **flowgraph-based serialisation library**
- A flowgraph is used to represent the Quartus **data dependency graph**, which when fed to a parallel computing library can yield huge reductions in loading time

## QUARTUS DATABASE SCHEMA PARSER V2

- Quartus' in-house **serialisation library** ships with a tool to automatically **generate serialisation routines (schemas)** based on C++ code
- This parser/generator tool featured a handwritten C++ parser, which was fragile and frequently broke whenever a new compiler was introduced
- **Redesigned and rewrote the parser/generator tool** to leverage **libclang** as its C++ parser
- This change improved the accuracy of the generated schemas while ensuring the long term viability of the library

## QUASI-SUCCINCT FIXED-DATA HASH TABLE

- Designed and implemented a **novel quasi-succinct hash table** data structure tailored for **precomputed data**
- Open-addressing & linear-probing based, where the worst case probe length was the size of the largest bucket
- **In memory compression** (inspired by **Elias-Fano encoding**) was used to store hash table indices. **x86 SIMD intrinsics** and jump tables were used to perform fast bucket offset lookup
- Invented as a solution to **eliminate wastage** when traditional associative data structures were used for **fixed data**

## HYBRID SHARED MUTEX & SEMI-LOCK FREE DATA STRUCTURES

- Designed to overcome the traditional limitations of conventional lock-free data structures, chiefly the need for individually-allocated nodes
- While various optimisations exist, node-based data structures generally have a high memory overhead
- **Hybrid shared mutex** leverages **thread-local-storage** to provide very **cheap read locks** due to its hardware cache advantages over shared atomics
- **Semi-lock free data structures** may use read locks to perform atomic operations on a **fixed size bufer**, while a write lock is used to increase the size of the buffer
- The ability to use **fixed size bufers** in semi-lock free data structures allows for more **compact implementations of high performance, thread-safe data structures**

## QUARTUS TCL VIRTUAL FILESYSTEM

- Implemented a **virtual filesystem layer** to provide the **Quartus Tcl interpreter** direct access to Quartus databases
- This greatly improves the project migration (export/import) experience by packaging all **user SDC and Tcl scripts** directly into the **Quartus databases**